

Block 4 – Cloud Computing

S2: Cloud Computing & Cloud Architecture

Arquitectura de Xarxes

Universitat Pompeu Fabra



- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing
- 3 Cloud & Container Platforms
- 4 Cloud Network Architecture
- 5 Cloud Network Services
- 6 Session Summary

Outline

- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing
- 3 Cloud & Container Platforms
- 4 Cloud Network Architecture
- 5 Cloud Network Services
- 6 Session Summary


If We Can Virtualize Everything... Why Not Sell It?

*Virtualization is the technology.
Cloud computing is the business model.*

Learning objectives:

- 1 Understand cloud computing's definition and essential characteristics
- 2 Distinguish IaaS (Infrastructure as a Service), PaaS (Platform as a Service), and SaaS (Software as a Service)
- 3 Understand core cloud architecture patterns: regions, availability zones, and virtual networks
- 4 Design a cloud virtual network with subnets, security controls, and network services

Recap – What We Virtualized (Session 1)



- **Compute:** VMs share physical hardware via hypervisors; containers share the host kernel directly
- **Networking:** Virtual NICs, Switches, and Routers replicate the physical stack in software
- **Isolation:** VLANs and overlay networks separate tenants on shared infrastructure

The Leap to Cloud

- Virtualization = technology (**how** you run multiple workloads)
- Cloud = business model (**what** you sell to customers)

Virtualization	Cloud Computing
Run multiple VMs/containers on one host	Offer compute as a service
Internal IT optimization	External/internal consumption
Manual provisioning possible	Automated, self-service
You own the hardware	Provider owns the hardware

Key difference: **automation and self-service** (no phone calls, no tickets, no waiting).

Discussion: From Virtualization to Cloud

*Your company already runs VMs on its own servers.
Why would you move to the cloud instead of
keeping everything in-house?*

Hints: think about CapEx vs OpEx, scaling speed, and who manages what.

Outline

- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing**
- 3 Cloud & Container Platforms
- 4 Cloud Network Architecture
- 5 Cloud Network Services
- 6 Session Summary

Cloud Computing – Definition (NIST)

*A model for enabling ubiquitous, convenient, **on-demand network access** to a shared pool of configurable computing resources that can be rapidly provisioned and released with minimal management effort.*

- **Ubiquitous:** anywhere, any device
- **On-demand:** no tickets, provision instantly
- **Shared pool:** many tenants, same hardware
- **Configurable:** choose CPU, RAM, storage, network
- **Rapidly provisioned:** up and running in seconds
- **Minimal management:** provider handles the hardware


Source: NIST SP 800-145 (2011). Still the standard definition used industry-wide.

Five Essential Characteristics

Characteristic	What it means
On-demand self-service	Provision resources via portal/API, no human interaction
Broad network access	Available from any device over the network
Resource pooling	Shared infrastructure, multi-tenant
Rapid elasticity	Scale up/down automatically (e.g. Black Friday → 10×)
Measured service	Pay only for what you use (CPU-hours, GB stored)

- **All five** must be present to qualify as “cloud computing”

Cloud Service Models – The Stack



IaaS, PaaS, SaaS – In Practice

IaaS (e.g. AWS EC2, Azure VMs, Google Compute Engine):

- Provider gives you VMs, storage, networks. You manage everything else
- “I want a Linux server in Frankfurt, 4 CPUs, 16 GB RAM, right now”

PaaS (e.g. Google App Engine, AWS Elastic Beanstalk, Heroku):

- Provider gives you runtime + tools. You only manage your code and data
- “I wrote a Python web app. Just run it, I do not care about servers”

SaaS (e.g. Google Workspace, Microsoft 365, Slack, Zoom):

- Provider gives you a complete application. You manage your data only
- “I need email for 500 employees. I do not want to run a mail server”

Service Models – Who Manages What?

Component	IaaS	PaaS	SaaS
Applications	You	You	Provider
Data	You	You	You
Runtime / Middleware	You	Provider	Provider
Operating System	You	Provider	Provider
Virtualization	Provider	Provider	Provider
Hardware	Provider	Provider	Provider

- IaaS → SaaS: **less control, less responsibility**
- Major providers (AWS, Azure, GCP) offer all three models

Discussion: Cloud Fundamentals

*A university wants to offer a web-based tool
for students to submit assignments.
Should they choose IaaS, PaaS, or SaaS? Why?*

Hints: consider the IT team size, maintenance burden, and how much control they need.

Outline

- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing
- 3 Cloud & Container Platforms**
- 4 Cloud Network Architecture
- 5 Cloud Network Services
- 6 Session Summary

Putting It All Together

- In Session 1 we saw the **building blocks**: Virtual Switches, Virtual Routers, Virtual NICs, VLANs
- These components do not run in isolation; they are managed by **platforms**
- Platforms orchestrate **VMs** or **containers** and configure virtual networking under the hood

Two categories:

- **Cloud platforms**: manage infrastructure at scale
 - Compute, storage, networking, and **VMs**
- **Container orchestration platforms**: deploy and scale **containerized** applications

Cloud Platforms

- **Self-managed / Private Cloud:** you own or rent the servers, install and manage the virtualization software yourself. Full control, but requires dedicated staff and expertise.

Platform	Description	Logo
OpenStack	Open-source, widely used in private clouds	 openstack®
VMware vSphere	Enterprise leader, proprietary	
Proxmox VE	Open-source, lightweight alternative	

- **Managed / Public Cloud:** rent capacity on demand, pay-as-you-go. Provider handles hardware, networking, cooling, security, and updates.

Provider	Description	Logo
Amazon Web Services	Market leader since 2006	
Microsoft Azure	Strong enterprise integration	
Google Cloud Platform	Data and AI focus	

These map directly to the IaaS model we just covered: the provider (or you) manages compute, storage, and networking.

Container Orchestration Platforms

- **Self-hosted / Private:** you install and operate on your own servers

Platform	Description	Logo
Docker	Container runtime; works well on a single host Managing hundreds of containers across hosts? Unmanageable	
Kubernetes (K8s)	Created to solve this: scheduling, scaling, networking, self-healing	 kubernetes
OpenShift	Enterprise K8s by Red Hat; adds security, CI/CD, UI	 Red Hat OpenShift

- **Hosted / Public:** cloud providers offer the same platforms as managed services. Each provider uses different names for essentially the same product (e.g., AWS ECS/EKS, Azure AKS, Google GKE, OpenShift on all major clouds).

Discussion: Platforms

A startup with 5 engineers needs to deploy a web application that may go from 100 to 100,000 users. Would you build a private cloud or use a public one? Why?

Hints: team size, upfront cost, time to market, scaling needs, control over infrastructure.

Outline

- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing
- 3 Cloud & Container Platforms
- 4 Cloud Network Architecture**
- 5 Cloud Network Services
- 6 Session Summary


Virtual Networks in the Cloud

- Every cloud provider lets you create your own **isolated virtual network**
 - AWS: Virtual Private Cloud (VPC)
 - Azure: Virtual Network (VNet)
 - GCP: VPC Network
- Think of it as your **private data center network** in the cloud
- You define:
 - **Address space** (e.g., 10.0.0.0/16)
 - **Subnets** (subdivisions of the address space)
 - **Routing rules** and **security policies**

Key: under the hood, cloud providers use **overlay networks** (covered in Block 5) to isolate tenants at scale.

Virtual Network Architecture

Virtual Network: 10.0.0.0/16



Subnets – Public vs Private

Cloud providers distinguish between subnets reachable from the Internet and isolated ones: not all workloads should be exposed. Web servers need public access, while databases and backends must stay hidden. Subnets enforce this boundary at the network level.

Public subnet:

- Has a route to an **Internet Gateway**
- Instances can have **public IP addresses**
- Use for: web servers, load balancers, bastion hosts

Private subnet:

- **No route** to the Internet; isolated by design
- Instances only have **private IPs**
- Use for: databases, internal services, backends

Route Tables

Every virtual network has an implicit virtual router managed by the provider. Route tables tell that router where to forward traffic.

- Each subnet is associated with a **route table**
- Route table = set of rules determining where traffic goes

Destination	Target / Next Hop	Subnet type
10.0.0.0/16	local (within virtual network)	Public + Private
0.0.0.0/0	Internet Gateway	Public only

- **Local route:** traffic within the virtual network stays inside (always present)
- **Default route** (0.0.0.0/0): public subnets exit via the Internet Gateway
- Private subnets have no default route; they are isolated by design

Internet Gateway

In a cloud virtual network, instances are isolated by default. The Internet Gateway is the controlled entry and exit point between a virtual network and the Internet.

- **Outbound only:** instance can reach the Internet, but is not reachable from outside
 - No public IP needed; the provider performs NAT transparently
- **Bidirectional:** instance is also reachable from the Internet
 - Requires a **public IP** assigned to the instance
 - Use for: web servers, load balancers, anything public-facing

Note: a NAT Gateway can optionally be added for private instances that need outbound Internet access (e.g. downloading updates), without exposing them inbound.

Security Groups – Instance-Level Firewall

Once traffic reaches an instance, you still need to control what it can send and receive. Security groups act as a per-instance firewall.

- **Security group** = virtual firewall attached to an instance
- Rules specify allowed **inbound** and **outbound** traffic
- **Stateful**: allow inbound → response automatically allowed

Direction	Protocol	Port	Source/Dest
Inbound	TCP	80	0.0.0.0/0
Inbound	TCP	443	0.0.0.0/0
Inbound	TCP	22	10.0.0.0/16
Outbound	All	All	0.0.0.0/0

Source: AWS, “Security Groups for Your VPC,” docs.aws.amazon.com.

Network ACLs (Access Control Lists) – Subnet-Level Firewall

Security groups protect individual instances. Network ACLs add an additional layer of control at the subnet boundary, before traffic even reaches any instance.

- **Network ACL** = firewall at the **subnet** level
- Rules evaluated in **order** (numbered, first match wins)
- **Stateless**: must explicitly allow both inbound AND outbound

Rule #	Direction	Protocol	Port	Action
100	Inbound	TCP	80	ALLOW
200	Inbound	TCP	443	ALLOW
*	Inbound	All	All	DENY

Source: AWS, “Network ACLs,” docs.aws.amazon.com.

Security Groups vs Network ACLs

	Security Groups	Network ACLs
Scope	Instance level	Subnet level
State	Stateful	Stateless
Rules	Allow only	Allow + Deny
Evaluation	All rules evaluated	First match wins
Default	Deny all inbound	Allow all

Defense in depth: use both together.

- Security Groups → fine-grained per-instance control
- Network ACLs → broad subnet-level guardrails

Discussion: Cloud Networking

*You deploy a web application in your cloud virtual network.
The web server needs to be public, but the database
must not be reachable from the Internet. How?*

Hints: think about public vs private subnets, security groups, and NAT gateways.

Outline

- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing
- 3 Cloud & Container Platforms
- 4 Cloud Network Architecture
- 5 Cloud Network Services**
- 6 Session Summary

Cloud Network Services – Overview

Beyond virtual networks and security controls, public cloud providers offer managed network services that simplify common infrastructure tasks.


- **Load Balancer (LB)**
- **Proxy**
- **Reverse Proxy**
- **VPN (Virtual Private Network)**

These services run as **managed offerings**: no servers to install or maintain.

Note: Proxy, Reverse Proxy, and VPN will be covered in depth in Block 6. Here we introduce them briefly.

Load Balancer

- Distributes incoming traffic across **multiple instances** (VMs or containers)
- Prevents any single instance from being overwhelmed
- If one instance fails, traffic is **redirected** to healthy ones




- **L4 (Transport Layer)** Load Balancer: routes based on IP and port (fast, simple)
- **L7 (Application Layer)** Load Balancer: routes based on HTTP headers, URL path, cookies (smarter)
- Cloud examples: AWS ALB/NLB, Azure Load Balancer, GCP Cloud Load Balancing

Source: AWS, "Elastic Load Balancing Documentation," docs.aws.amazon.com.

Proxy

Without a proxy, every instance in your network reaches the Internet directly: no visibility, no control. A proxy sits in front of clients and mediates all outbound traffic.

- Clients send requests to the proxy, which forwards them to the Internet




- **Access control:** block certain websites or domains
- **Caching:** store frequently accessed content, reduce bandwidth
- **Logging:** monitor what employees or VMs access
- Examples: Squid, corporate firewalls with proxy mode

Reverse Proxy

Without a reverse proxy, clients connect directly to your backend servers: exposing their addresses and requiring each to handle TLS. A reverse proxy sits in front of servers and mediates all inbound traffic.

- Clients never communicate directly with the backend



- **TLS termination:** handles HTTPS encryption, backends receive plain HTTP
- **URL routing:** /api/* → API server, /* → web app
- **Caching:** stores responses to reduce backend load
- **Examples:** NGINX, HAProxy, AWS CloudFront, Azure Front Door

Reverse Proxy vs Proxy

	Proxy	Reverse Proxy
Protects	Clients (outbound)	Servers (inbound)
Position	In front of clients	In front of servers
Who knows?	Server sees proxy, not client	Client sees proxy, not server
Use case	Control outbound access	TLS termination, routing, caching

Think of it this way:

- **Proxy:** “I control what my users can access outside”
- **Reverse proxy:** “I control how outside users reach my servers”

VPN – Virtual Private Network

- A **VPN (Virtual Private Network)** creates a **secure, encrypted tunnel** over the public Internet
- Connects remote networks or users as if they were on the same private network

VPNs will be covered in depth in Block 6.

Discussion: Cloud Network Services

*You are deploying a web application.
When does it make sense to add a Load Balancer?
When would you not need one?*

Hints: think about the number of users, availability requirements, and what happens if your single instance goes down.

Outline

- 1 From Virtualization to Cloud
- 2 Fundamentals of Cloud Computing
- 3 Cloud & Container Platforms
- 4 Cloud Network Architecture
- 5 Cloud Network Services
- 6 Session Summary**

Key Takeaways

- 1 Virtualization is the **technology**; cloud is the **business model**
- 2 Five NIST characteristics define true cloud computing
- 3 **IaaS / PaaS / SaaS** differ in who manages what
- 4 Cloud and container **platforms** deliver these models at scale
- 5 A **virtual network** is your isolated cloud network (CIDR, subnets, route tables)
- 6 **Security groups** (instance) + **ACLs** (subnet) = defense in depth
- 7 **Load balancers, reverse/forward proxies, and VPNs** are essential cloud network services

Discussion

*You are designing a cloud network for a company with public web servers and private databases.
Which components would you use and why?*

Think about: public vs private subnets, security groups, NAT gateways, load balancers, VPN.

References

- ❶ NIST SP 800-145, Mell & Grance, "The NIST Definition of Cloud Computing," 2011.
- ❷ NIST SP 800-144, "Guidelines on Security and Privacy in Public Cloud Computing," 2011.
- ❸ AWS, "Amazon VPC User Guide," docs.aws.amazon.com.
- ❹ Microsoft Azure, "Azure Virtual Network Documentation," docs.microsoft.com.
- ❺ Google Cloud, "VPC Documentation," cloud.google.com/vpc/docs.
- ❻ CSA, "Security Guidance for Critical Areas of Focus in Cloud Computing v5," 2022.
- ❼ Erl, Puttini & Mahmood, *Cloud Computing: Concepts, Technology & Architecture*, Prentice Hall, 2013.
- ❽ Kurose & Ross, *Computer Networking: A Top-Down Approach*, 8th ed., Pearson, 2021.
- ❾ AWS, "Elastic Load Balancing Documentation," docs.aws.amazon.com.
- ❿ NGINX, Inc., "What Is a Reverse Proxy?," nginx.com/resources.
- ⓫ AWS, "AWS VPN Documentation," docs.aws.amazon.com.